Data processing: Pictures are rescaled to 224 x 224; Data augmentation on train set (rescale, rotation, shift, shear, zoom, horizontal flip)
CNN architecture: Use the ImageNet pre-trained VGG16 as feature extractor (remove the top layer), We add a dense layer with 512 unit with dropout 0.5 and a dense layer with 4 units and softmaxloss function on top of conv_base.
Fine tuning: Freeze all the layers of the VGGnet before block5_conv1
Batch size (for the training/validation generators) = 32
Optimizer RMSprop with **learning rate** = 1e-5
Epochs = 20
The train and validation process is as following. And the accuracy on test dataset is 0.9850.

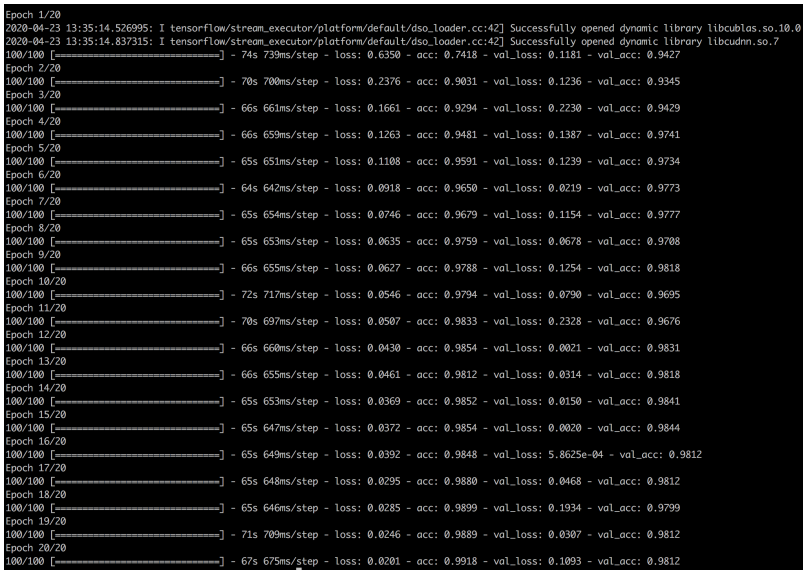

(Failed exercise) What's more, I tried to use ImageNet pre-trained Resnet50 with fine-tuning before. The result is as following.

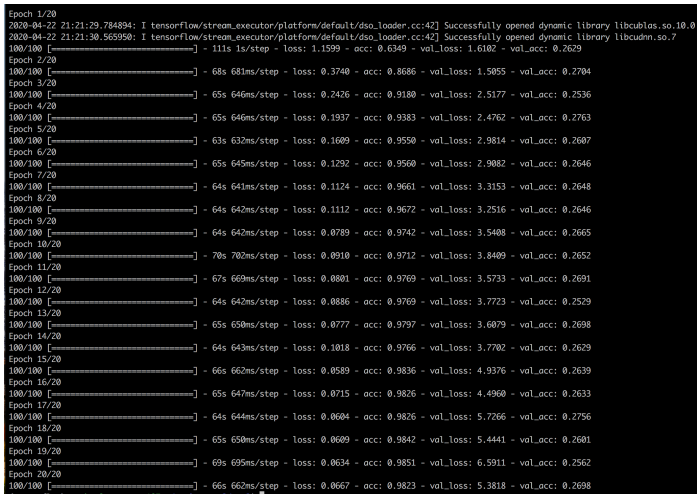


```
c2/car_0766.jpg: "cat"
c2/car_0767.jpg: "cat"
c2/car_0768.jpg: "cat"
c2/car_0769.jpg: "cat"
c2/car_0770.jpg: "cat"
c2/car_0771.jpg: "cat"
c2/car_0772.jpg: "cat"
c2/car_0773.jpg: "cat"
c2/car_0774.jpg: "cat"
c2/car_0775.jpg: "cat"
c2/car_0776.jpg: "cat"
c3/motorbike_0592.jpg: "cat"
c3/motorbike_0593.jpg: "cat"
c3/motorbike_0594.jpg: "cat"
c3/motorbike_0595.jpg: "cat"
c3/motorbike_0596.jpg: "cat"
c3/motorbike_0597.jpg: "cat"
c3/motorbike_0598.jpg: "cat"
```

The train loss keeps decreasing while the validation loss keeps increasing, which indicates an overfitting situation. The test result of this model also proves that because this model recognized everything as the cat label.
After all of this, I find VGG structure is good enough for this small dataset.